

Estimating Energy Cost

GSOC Mid-Term Evaluation

Manas Pratim Biswas

<https://github.com/sanam2405/SoftwareEnergyCost>

Progress Till Now :

1. Profiled Baler for **CPU computation** time with **cProfile** and Visualised the results with **SnakeViz** and **yelp-gprof2dot** (**Deterministic Profiling**)
2. Profiled Baler for **CPU computation** on wall clock time (**Sampling Profiling**)
3. Profiled Baler for **memory consumption** with **memory-profiler**
4. Profiled Baler for **power consumption** with **powermetrics** and Visualised results using **influxdb**
5. Profiled the **CO2 Emissions** and **power consumption** with **codecarbon**

cProfile

Deterministic Profiler. Traces each function call and measures CPU computation.

tottime - Total time taken by the method excluding recursive calls

cumtime - Cumulative Time (Time taken by the method including recursive calls)

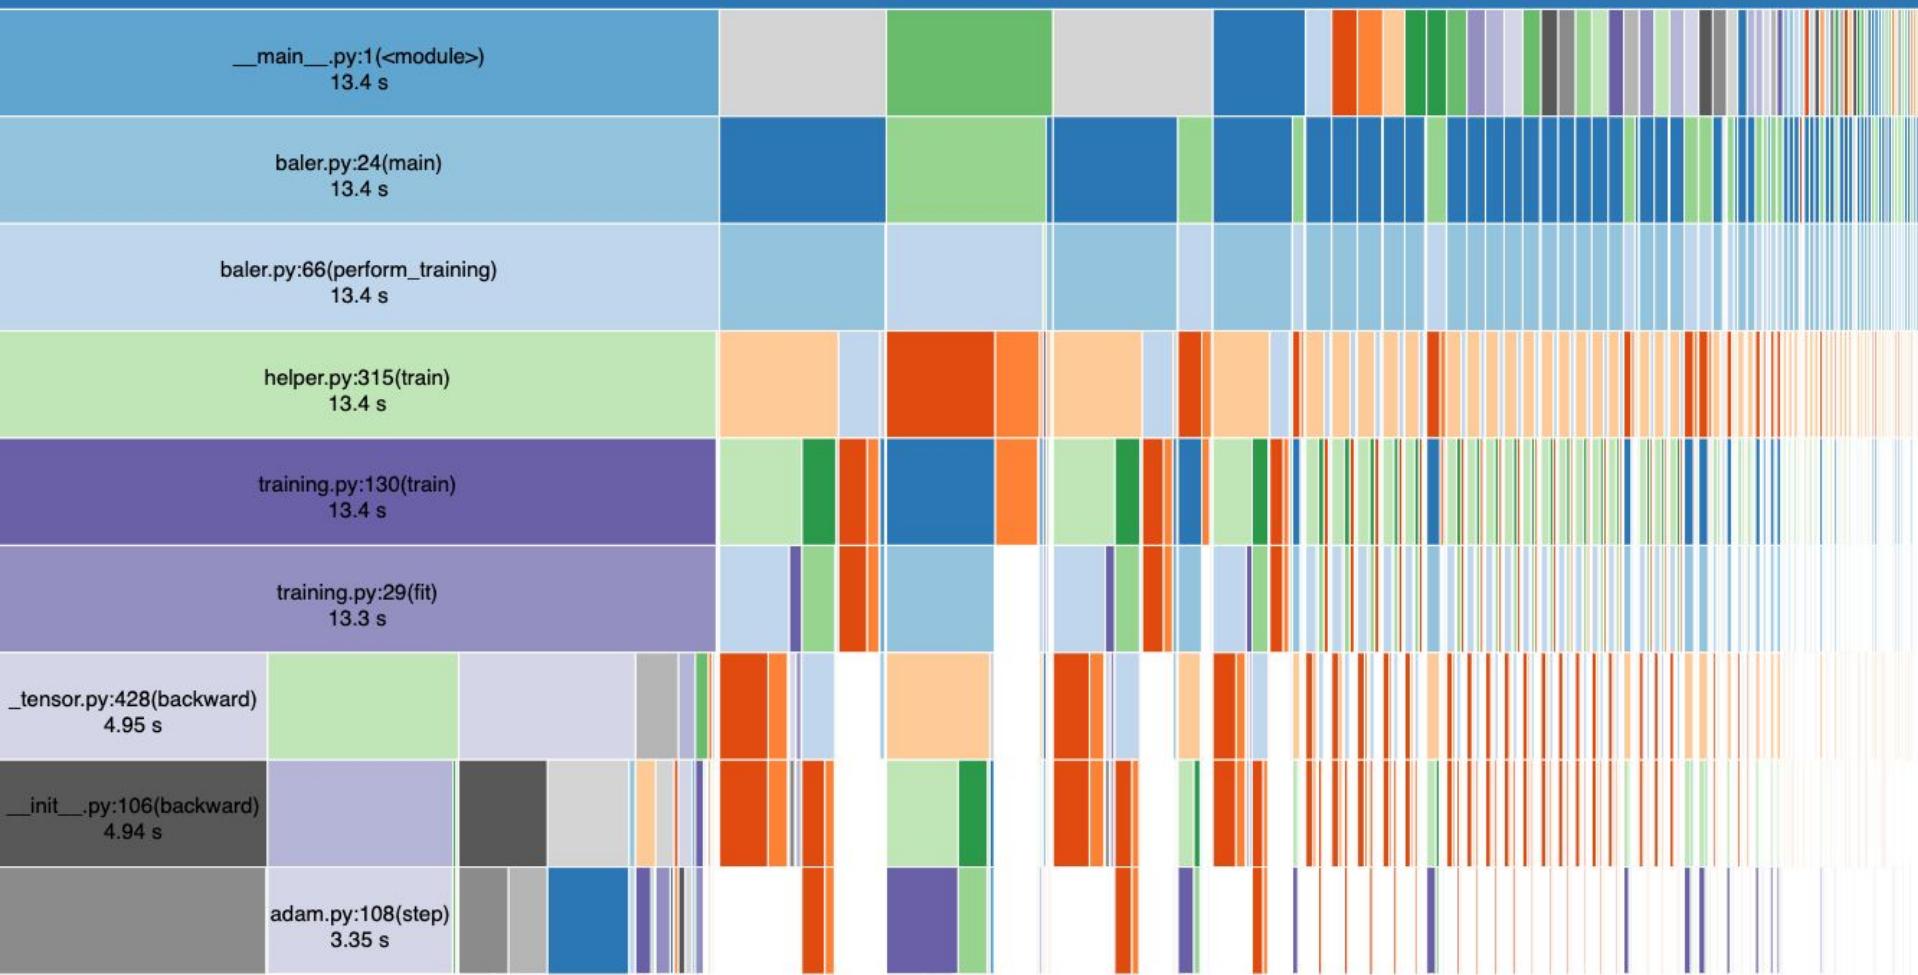
ncalls - Total number of times the method is invoked (Includes recursive call)

pcall - Total number of primitive calls to the method (Excludes recursive call)

All the results are calculate for 1000 epochs of train, compress and decompress on the CFD dataset. It takes **14.56 seconds** to run baler on this configuration on my Apple M1 silicon.

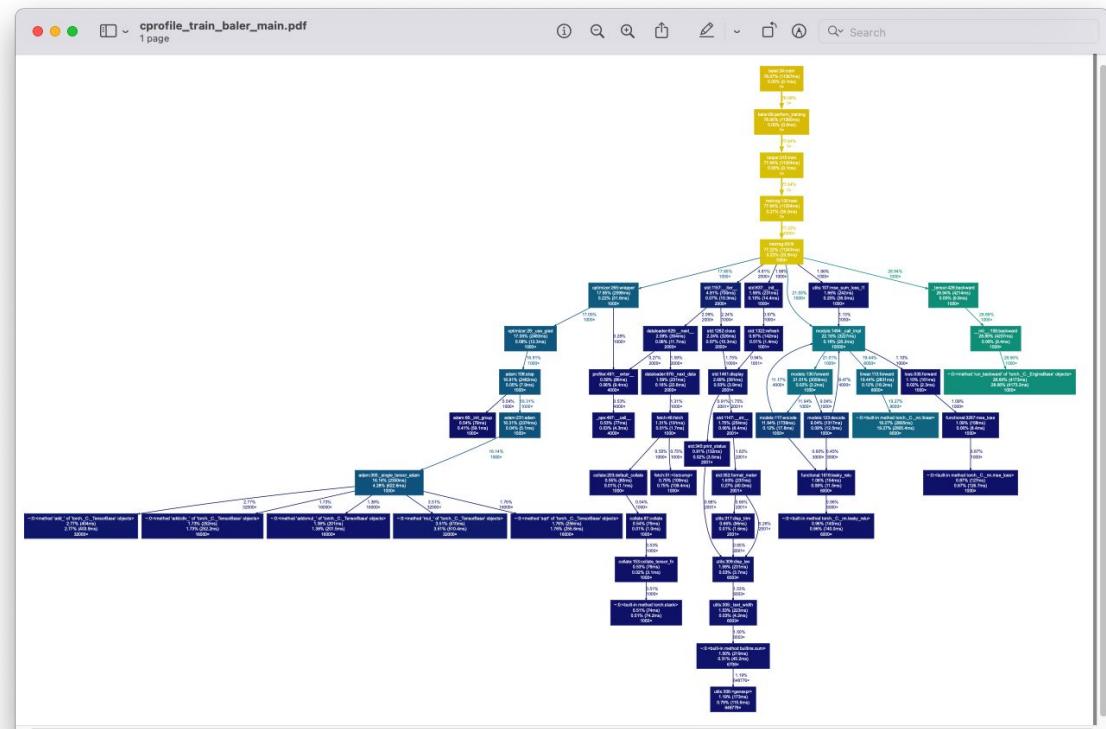
~:0(<built-in method builtins.exec>)

16.5 s



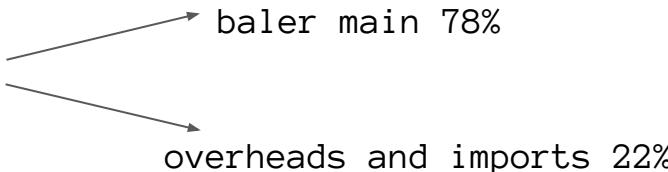
cProfile

- DI Graph generated while Training Baler
- Provides the exact call order for the methods
- Shows an exact breakdown of time distribution across various parts of the code



Execution of Baler(train)

Total time taken 14.56 second



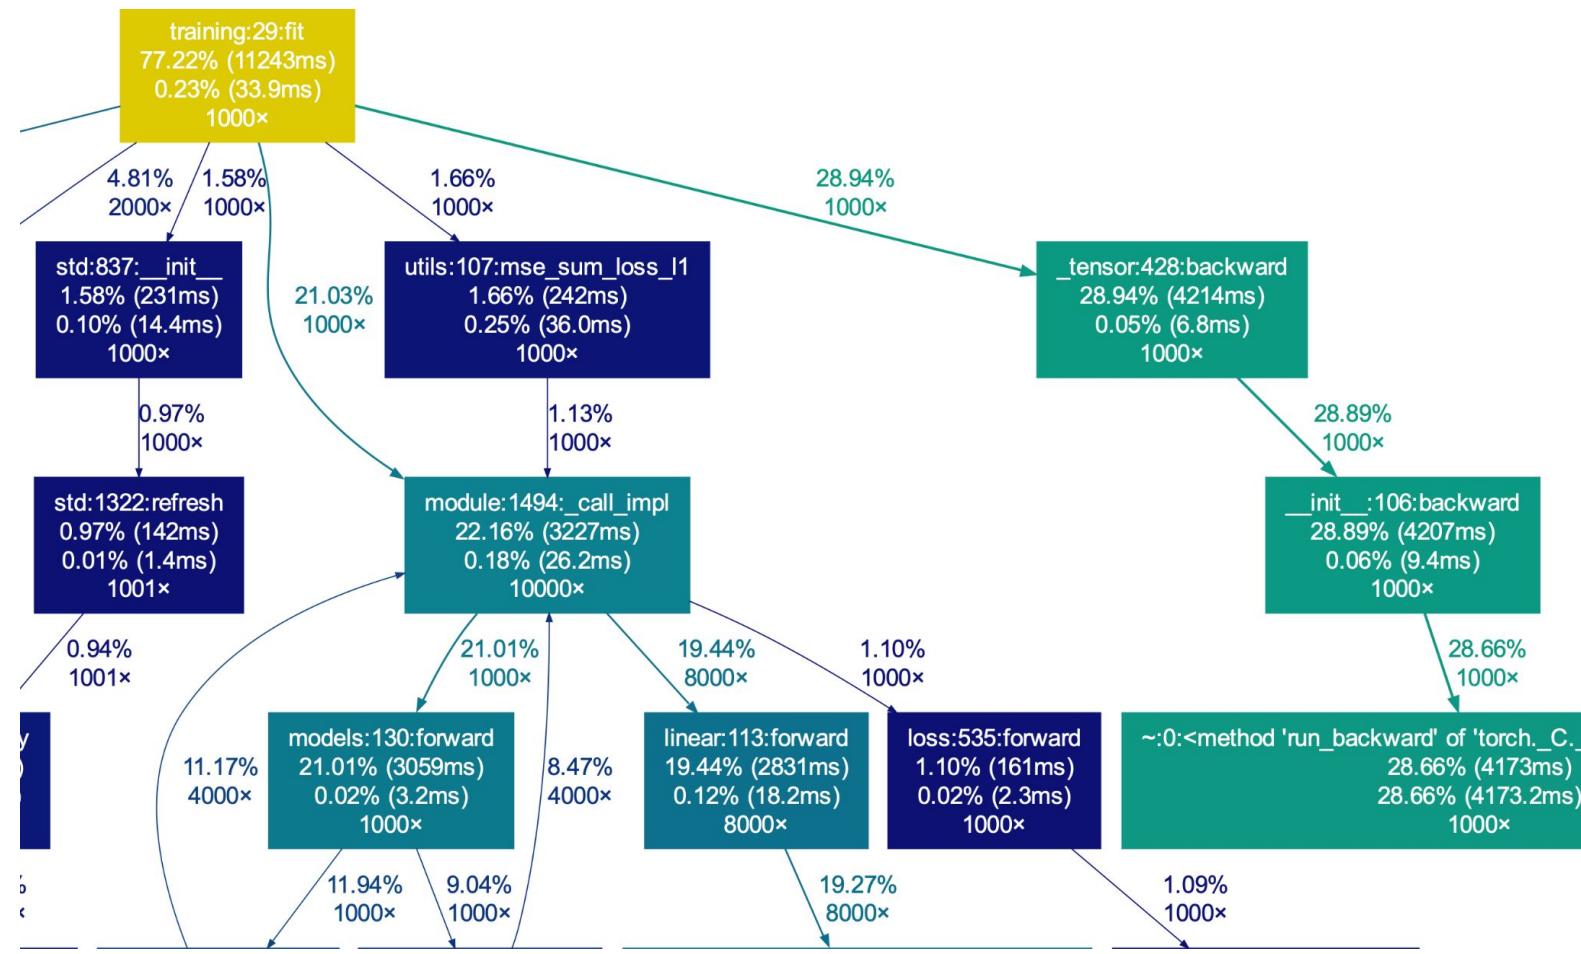
Our subject of interests are

- `train(...)` in `training.py`
- `fit(...)` in `training.py`

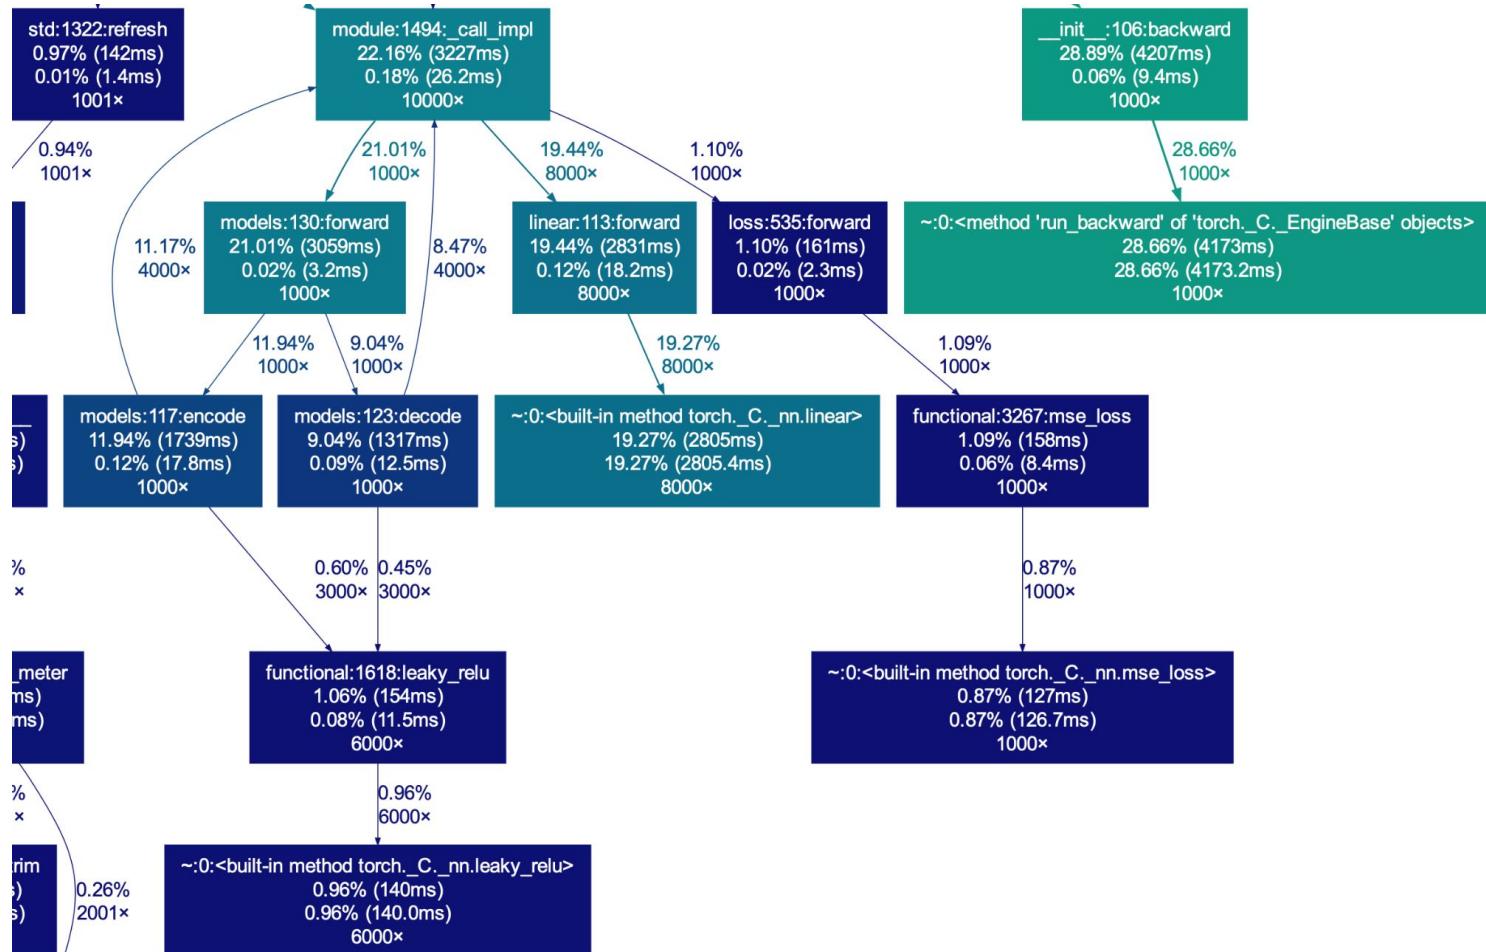


Train fit(...)

T(Backpropagation) > T(Forward)



Train fit(. . .) | T(encode) > T(decode)



cProfile (Deterministic)

Disadvantages

1. High overhead and distorted results
2. Too much information is given as tracing all the method invokes
3. Suitable to run and test on local environment. Doesn't consider production env
4. Profiles on the basis of method execution time and not on input or logic

Pyinstrument (Sampling)

1. Solves the issue of high overhead and distorted results
2. Concise information is given
3. Measures wall clock time contrary to cProfile which measured CPU time

Pyinstrument (Sampling Profiler)

Methodology

- Draws samples at some defined intervals of time
- Observe and make estimations based on the sampled data

Benefit

- A more statistically suited sampling method is used
- This uniforms the biases and outliers

Pyinstrument (Sampling Profiler)

Pyinstrument: How it works, cat edition

- 12:00: Sleeping 
- 12:05: Sleeping 
- 12:10: Eating 
- 12:15: Using the litterbox 
- 12:20: Sleeping 
- 12:25: Sleeping 
- 12:30: Sleeping 

Pyinstrument (Sampling Profiler)

Pyinstrument: How it works, cat edition

- Breakdown of observed cat activities
 - 90%: Sleeping 
 - 10%: Eating 
 - 9%: Using the litterbox 
 - 1%: Staring longingly through the window at birds


Wall Clock Time vs CPU Time

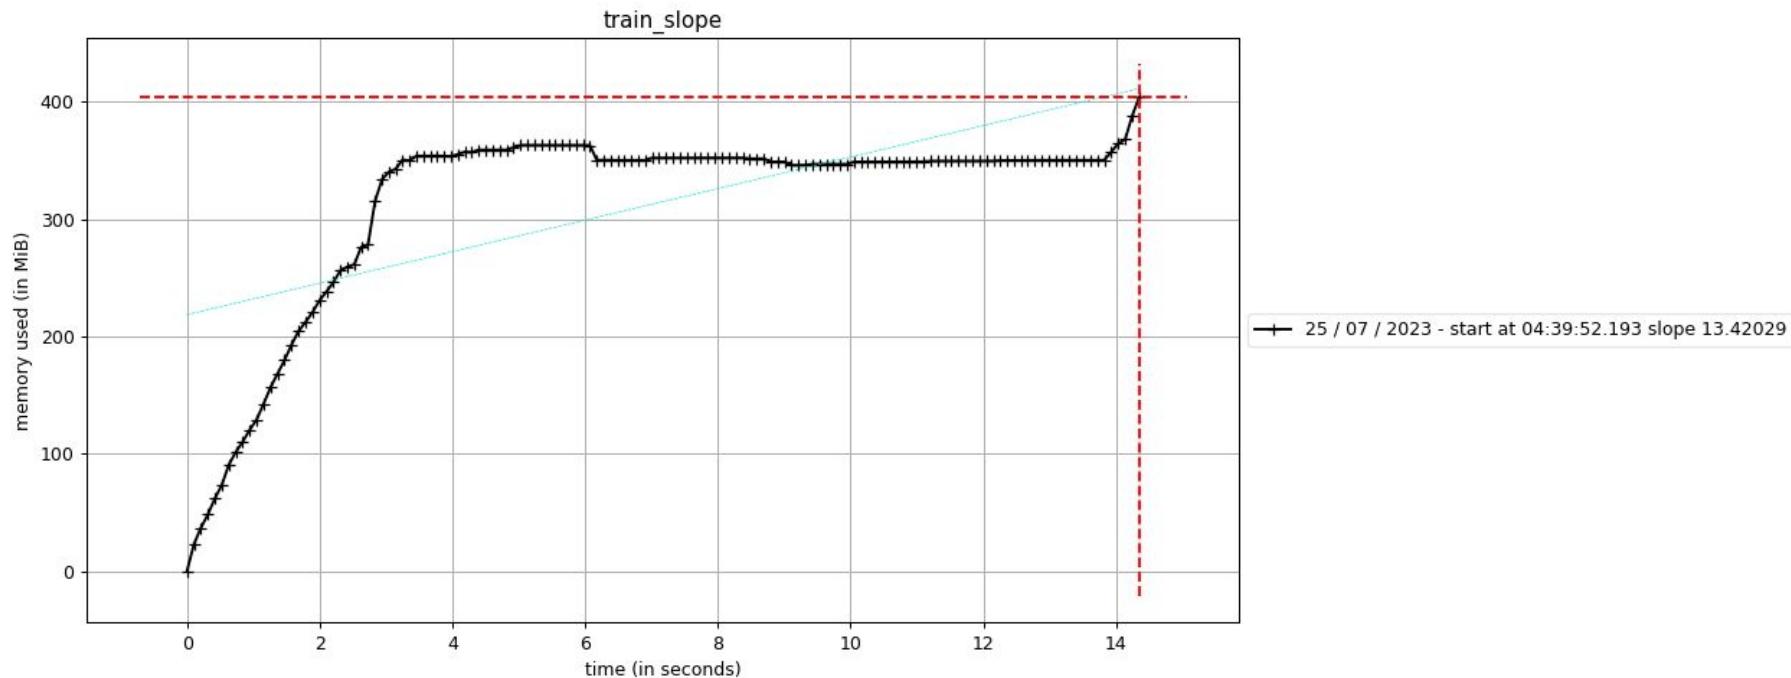
- Wall-clock time is the time that a clock on the wall (or a stopwatch in hand) would measure as having elapsed between the start of the process and 'now'
- The CPU time is the amount of time spent by the user code in the CPU
- The wall-clock time is hence not the number of seconds that the process has spent on the CPU rather it is the elapsed time, including time spent waiting for its turn on the CPU (while other processes get to run)
- Multicore processors, which incorporate multiple processing units, can run several threads at once to use more CPU time than wall clock time

Execution of baler(train)

```
██████████ Recorded: 19:42:01 Samples: 7999  
██████████ Duration: 14.427 CPU time: 82.388  
██████████ v4.5.0  
  
Program: baler --mode train --project CFD_workspace CFD_project_animation  
  
11.498 _run_code numpy.py:64  
└─ 11.498 <module> baler/_main__.py:1  
  └─ 11.498 main baler/baler.py:24  
    └─ 11.495 perform_training baler/baler.py:66  
      └─ 11.439 train baler/modules/helper.py:315  
        └─ 11.439 train baler/modules/training.py:130  
          └─ 11.392 fit baler/modules/training.py:29  
            └─ 4.384 Tensor.backward torch/_tensor.py:428  
              [7 frames hidden] torch, <built-in>  
              4.383 __EngineBase.run_backward None  
            └─ 3.613 CFD_dense_AE._call_impl torch.nn.modules.module.py:1494  
              [4 frames hidden] torch, <built-in>  
              3.610 CFD_dense_AE.forward baler/modules/models.py:130  
                └─ 1.954 CFD_dense_AE.encode baler/modules/models.py:117  
                  └─ 1.946 Linear._call_impl torch.nn.modules.module.py:1494  
                    [5 frames hidden] torch, <built-in>  
                    1.655 CFD_dense_AE.decode baler/modules/models.py:123  
                      └─ 1.638 Linear._call_impl torch.nn.modules.module.py:1494  
                        [10 frames hidden] torch, <built-in>  
            └─ 2.242 wrapper torch/optim/optimizer.py:265  
              [56 frames hidden] torch, <built-in>  
            └─ 0.827 tqdm.__iter__ tqdm/std.py:1157  
              [219 frames hidden] tqdm, <built-in>, _weakrefset, torch,...  
            └─ 0.158 tqdm.__init__ tqdm/std.py:837  
              [122 frames hidden] tqdm, <built-in>, _weakrefset, torch  
  
To view this report with different options, run:  
  pyinstrument --load-prev 2023-07-07T19-42-01 [options]  
  
(base) manas@Manass-MacBook-Air baler %
```

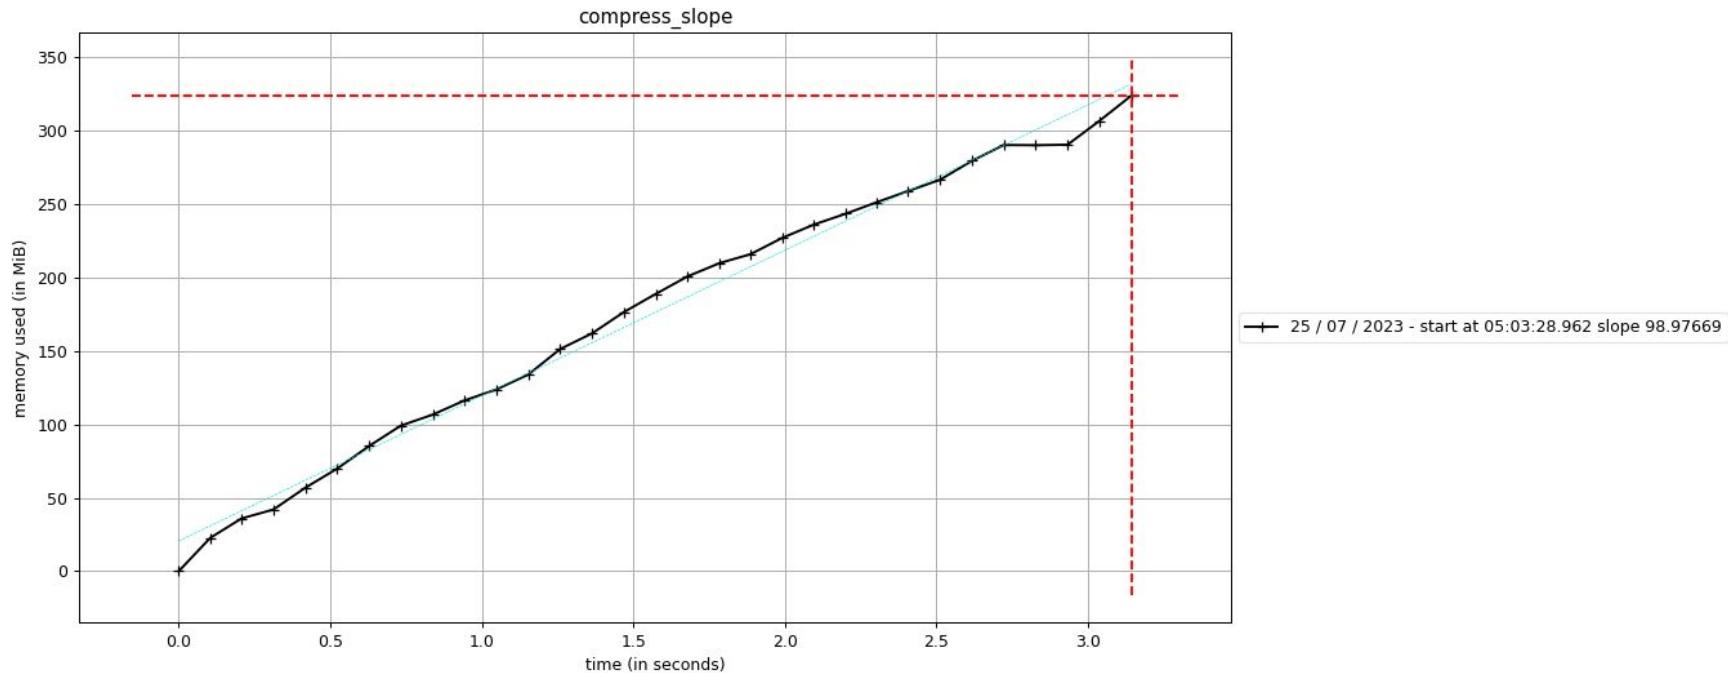
Memory-profiler (MiB – mebibyte)

Train 1000 epochs



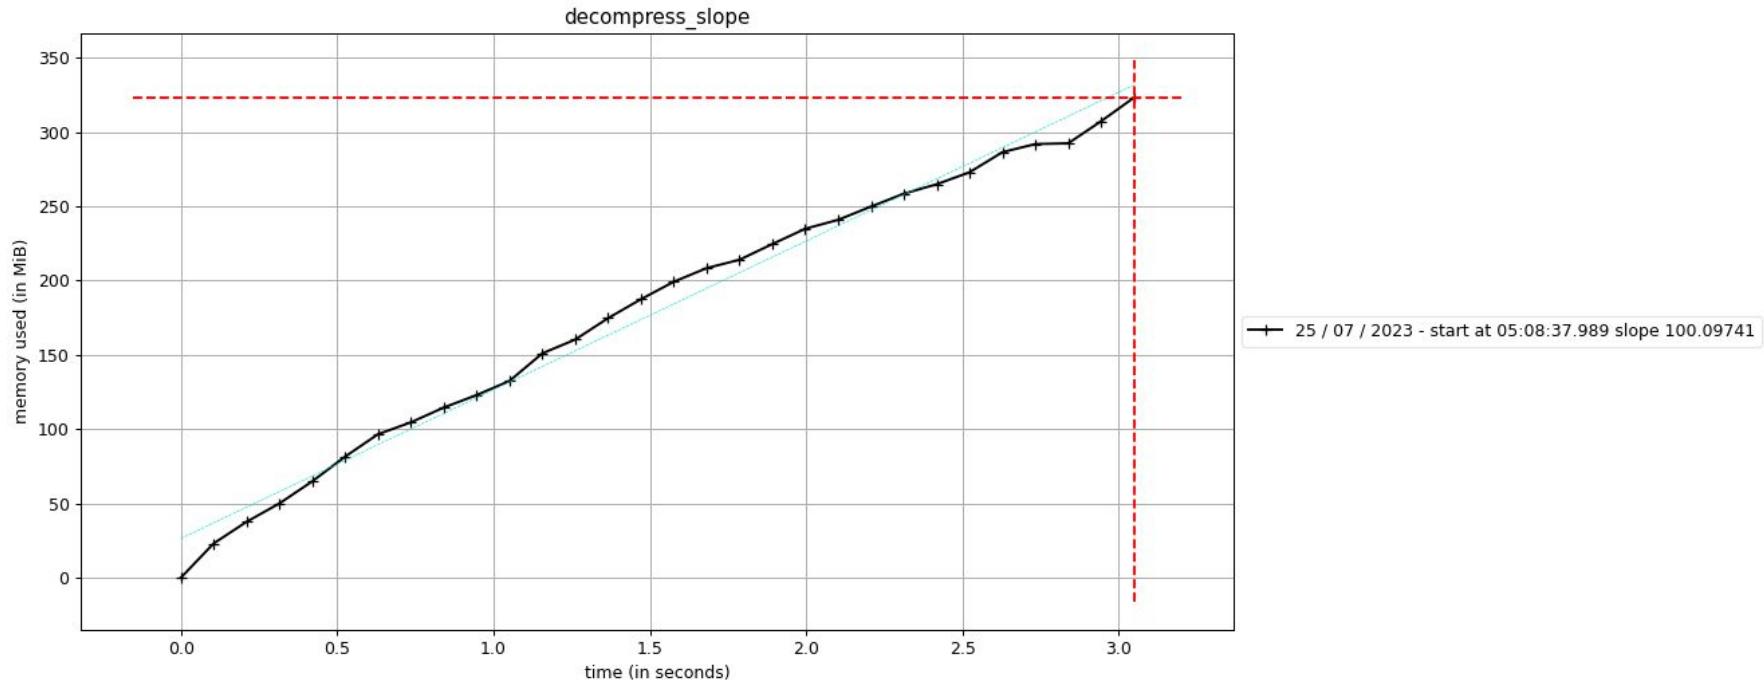
Memory-profiler (MiB – mebibyte)

Compress 1000 epochs



Memory-profiler (MiB – mebibyte)

Decompress 1000 epochs



codecarbon

Methodology – <https://mlco2.github.io/codecarbon/methodology.html>

- C = Carbon Intensity of the electricity consumed for computation: quantified as g of CO₂ emitted per kilowatt-hour of electricity.
- P = Power Consumed by the computational infrastructure: quantified as kilowatt-hours.
- Carbon dioxide emissions (CO₂eq) can then be calculated as C * P

Carbon Intensity of the consumed electricity is calculated as a weighted average of the emissions from the different energy sources that are used to generate electricity, including fossil fuels and renewables.

A scaling factor of **1e6** was been used to generate the plots for **50 baler runs with 1000 epochs each**

Data Fields Logged for Each Experiment

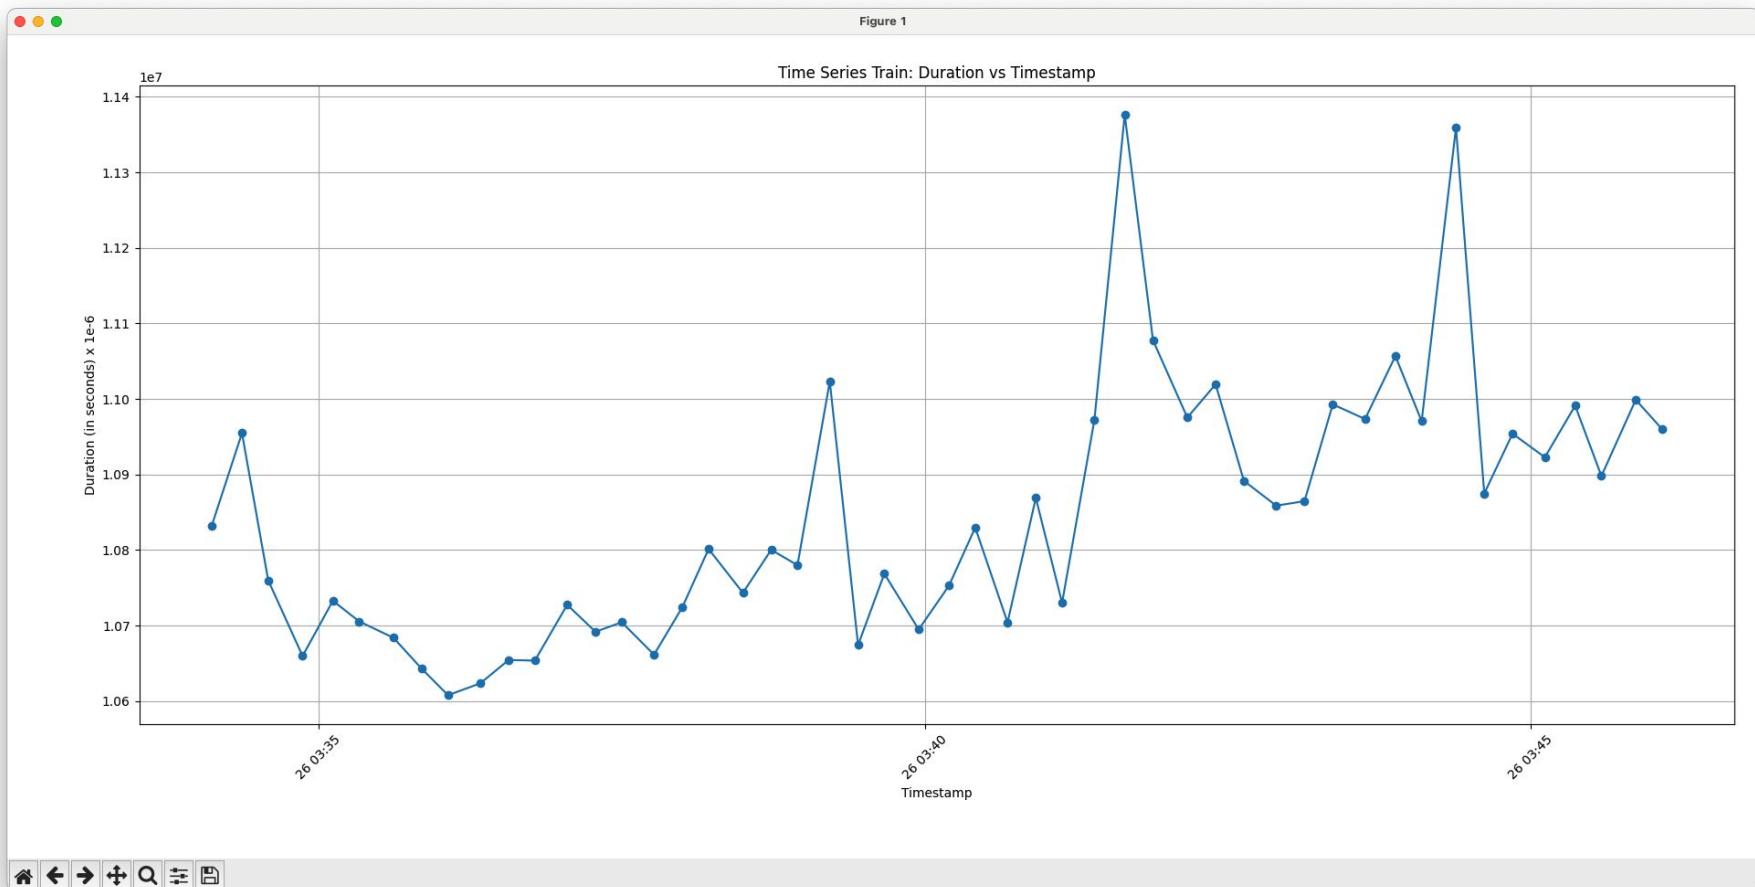
Field	Description
timestamp	Time of the experiment in <code>%Y-%m-%dT%H:%M:%S</code> format
project_name	Name of the project, defaults to <code>codecarbon</code>
run-id	id of the run
duration	Duration of the compute, in seconds
emissions	Emissions as CO ₂ -equivalents [CO ₂ eq], in kg
emissions_rate	emissions divided per duration, in Kg/s
cpu_power	CPU power (W)
gpu_power	GPU power (W)
ram_power	RAM power (W)
cpu_energy	Energy used per CPU (kW)
gpu_energy	Energy used per GPU (kW)
ram_energy	Energy used per RAM (kW)
energy_consumed	sum of cpu_energy, gpu_energy and ram_energy (kW)
country_name	Name of the country where the infrastructure is hosted
country_iso_code	3-letter alphabet ISO Code of the respective country
region	Province/State/City where the compute infrastructure is hosted

codecarbon

Synopsis

- CO2 Emission
 - $0.58 * 1e-6$ kg CO2 eqv while train
 - $0.44 * 1e-6$ kg CO2 eqv while compress
 - $0.59 * 1e-6$ kg CO2 eqv while decompress
- Energy
 - $15.35 * 1e-6$ kWh energy per consumed while train
 - $0.078 * 1e-6$ kWh energy per consumed while compress
 - $0.064 * 1e-6$ kWh energy per consumed while decompress

Figure 1



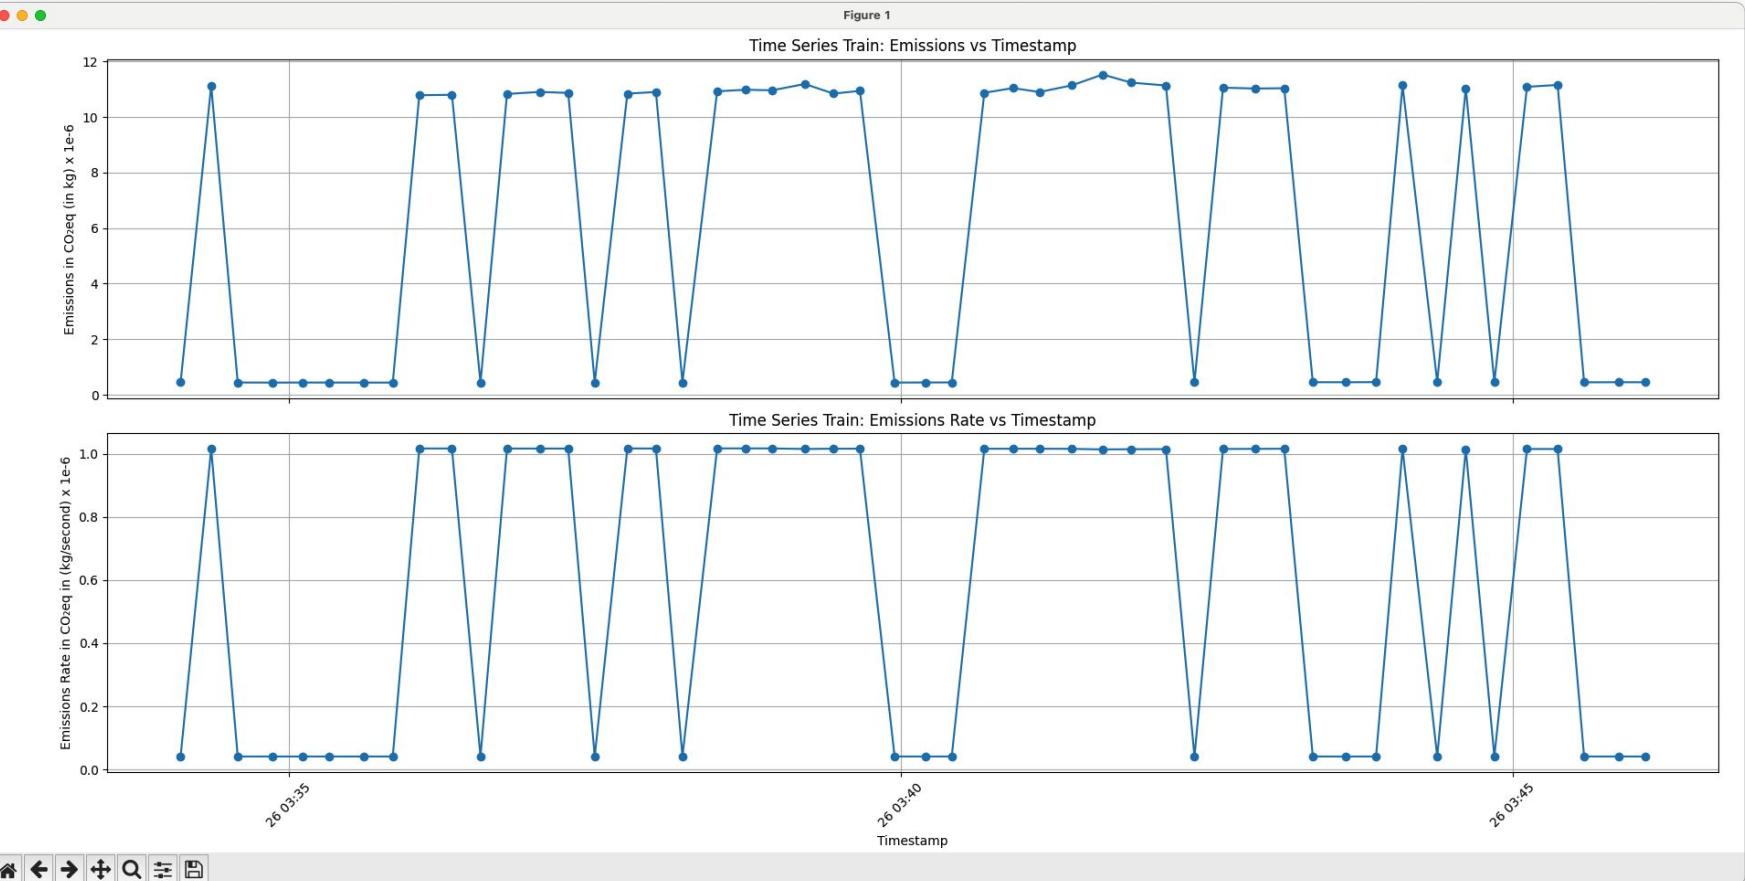


Figure 1

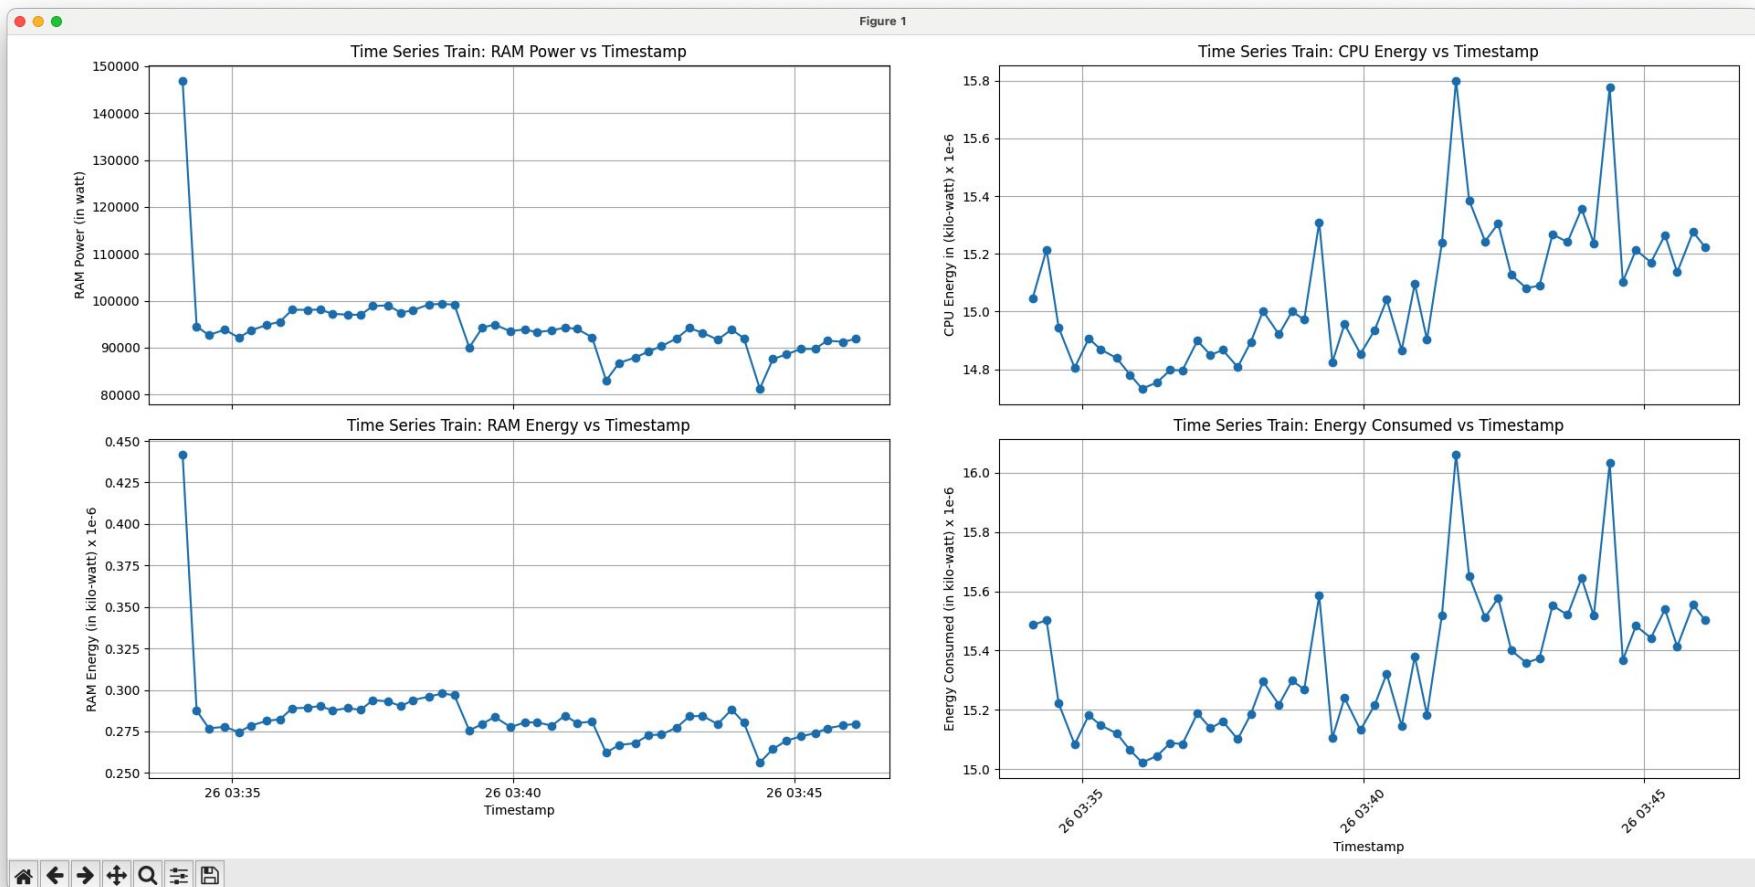


Figure 1

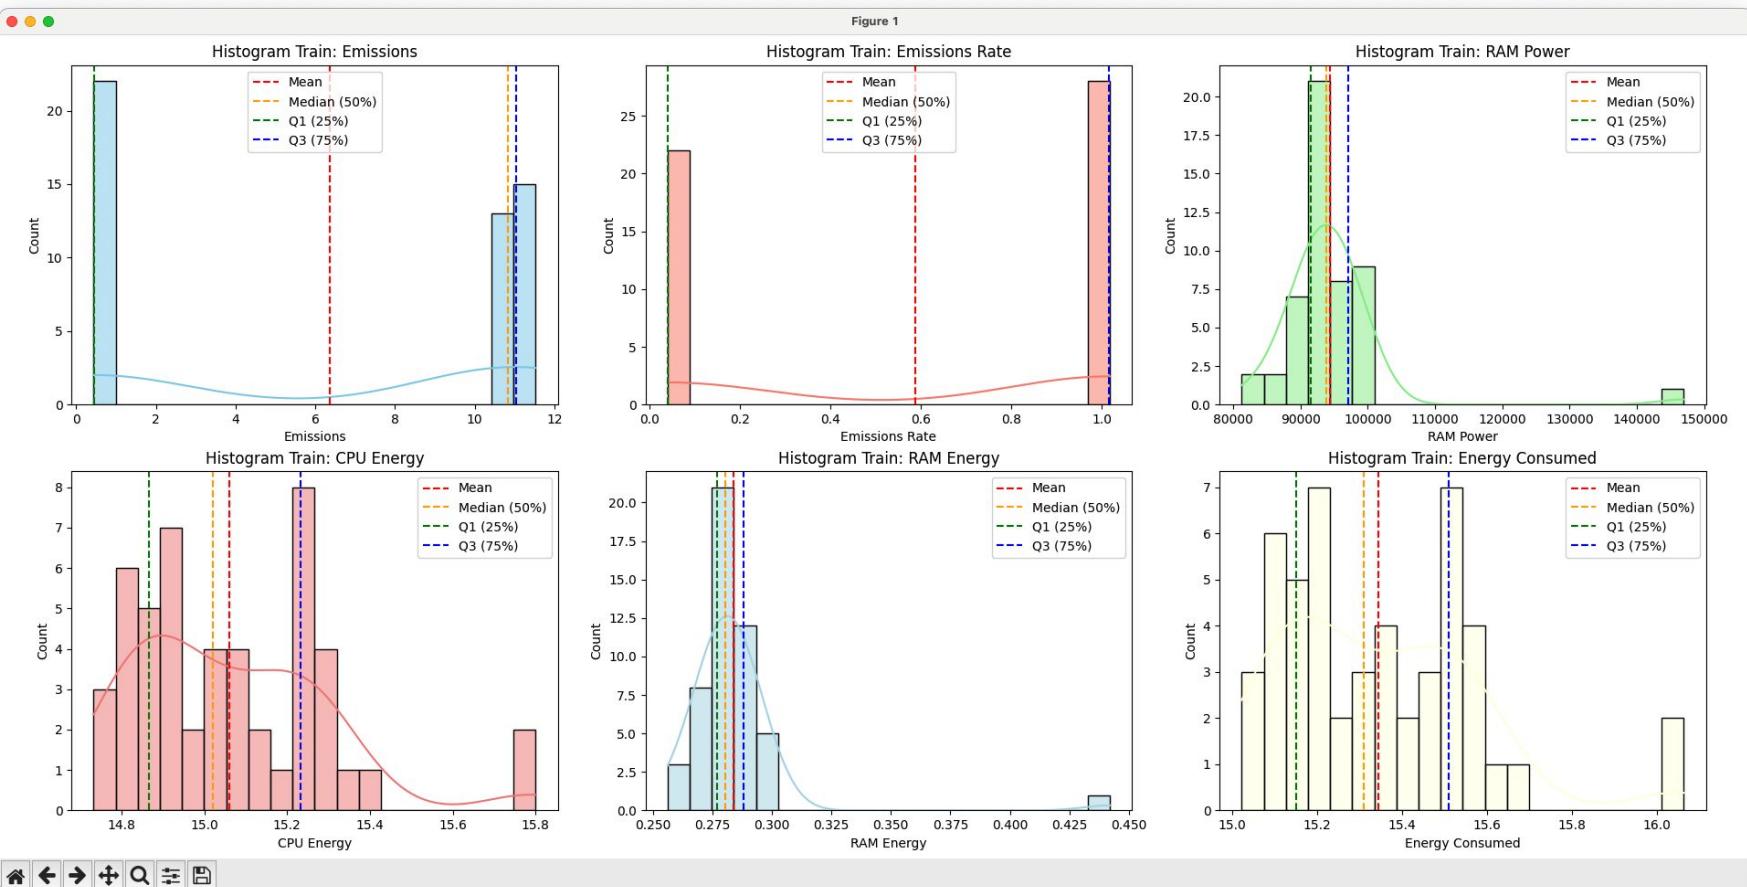


Figure 1

Time Series Compress: Duration vs Timestamp

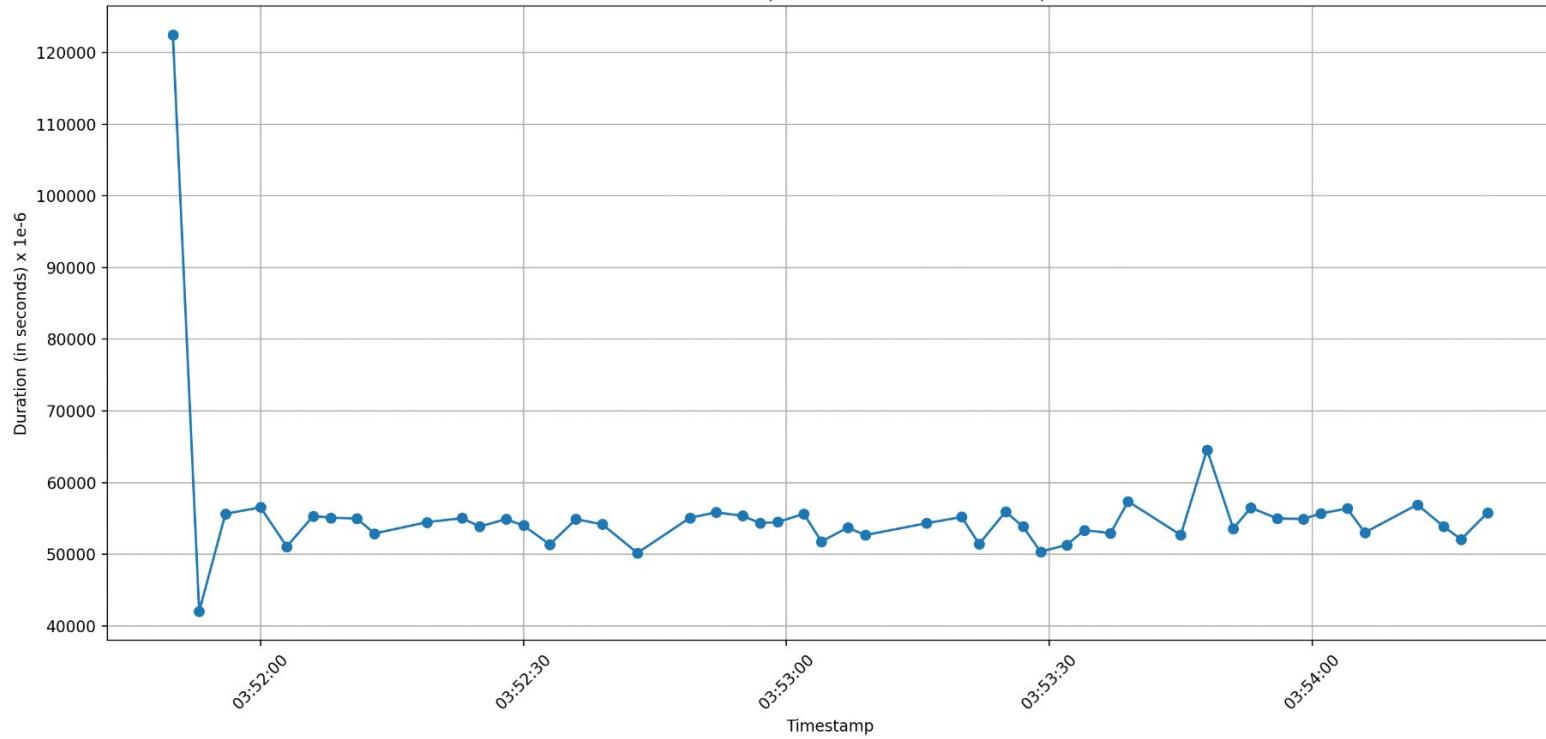


Figure 1

Time Series Compress: Emissions vs Timestamp

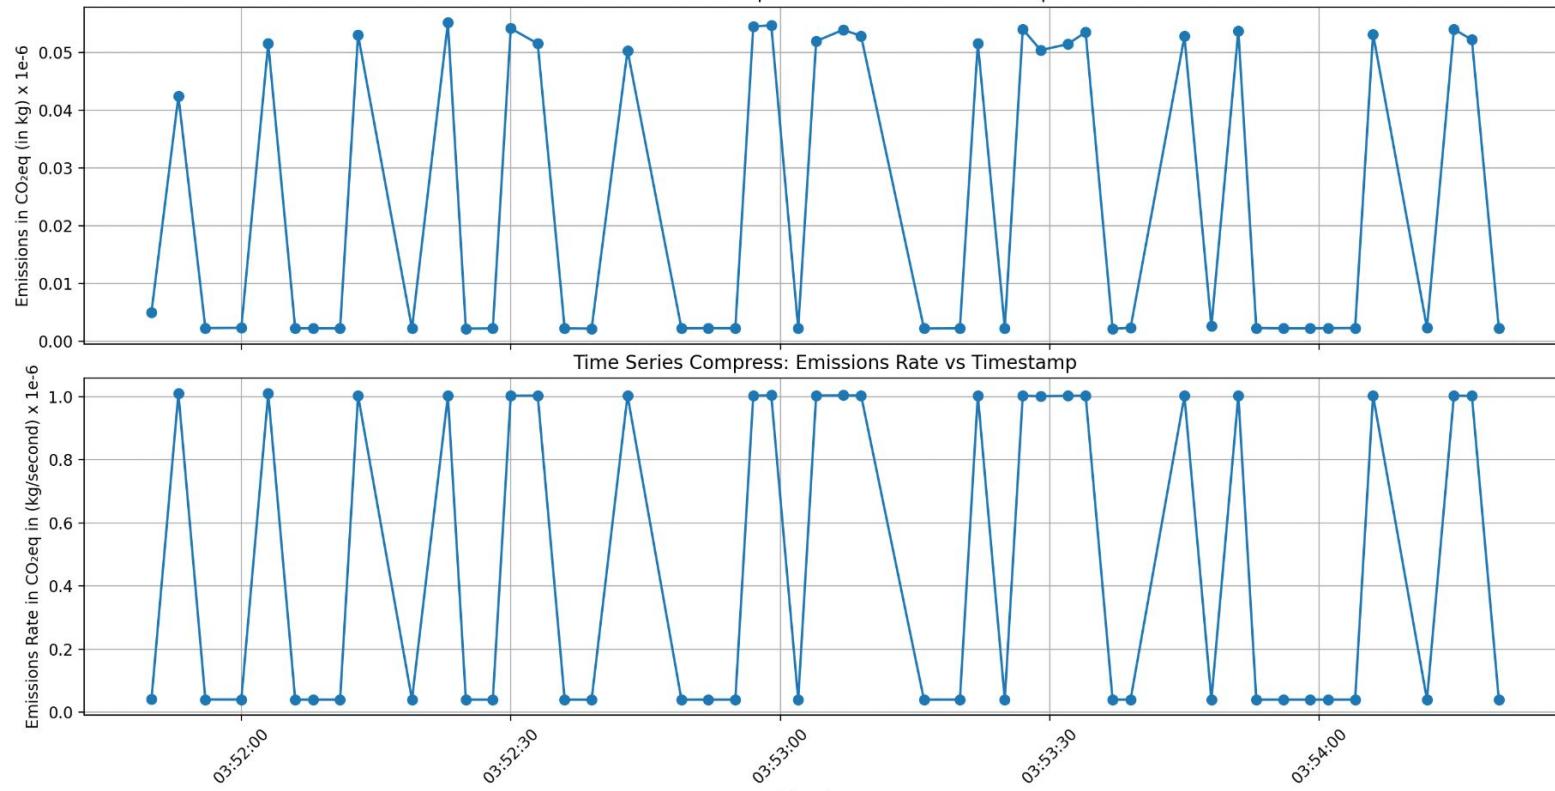


Figure 1

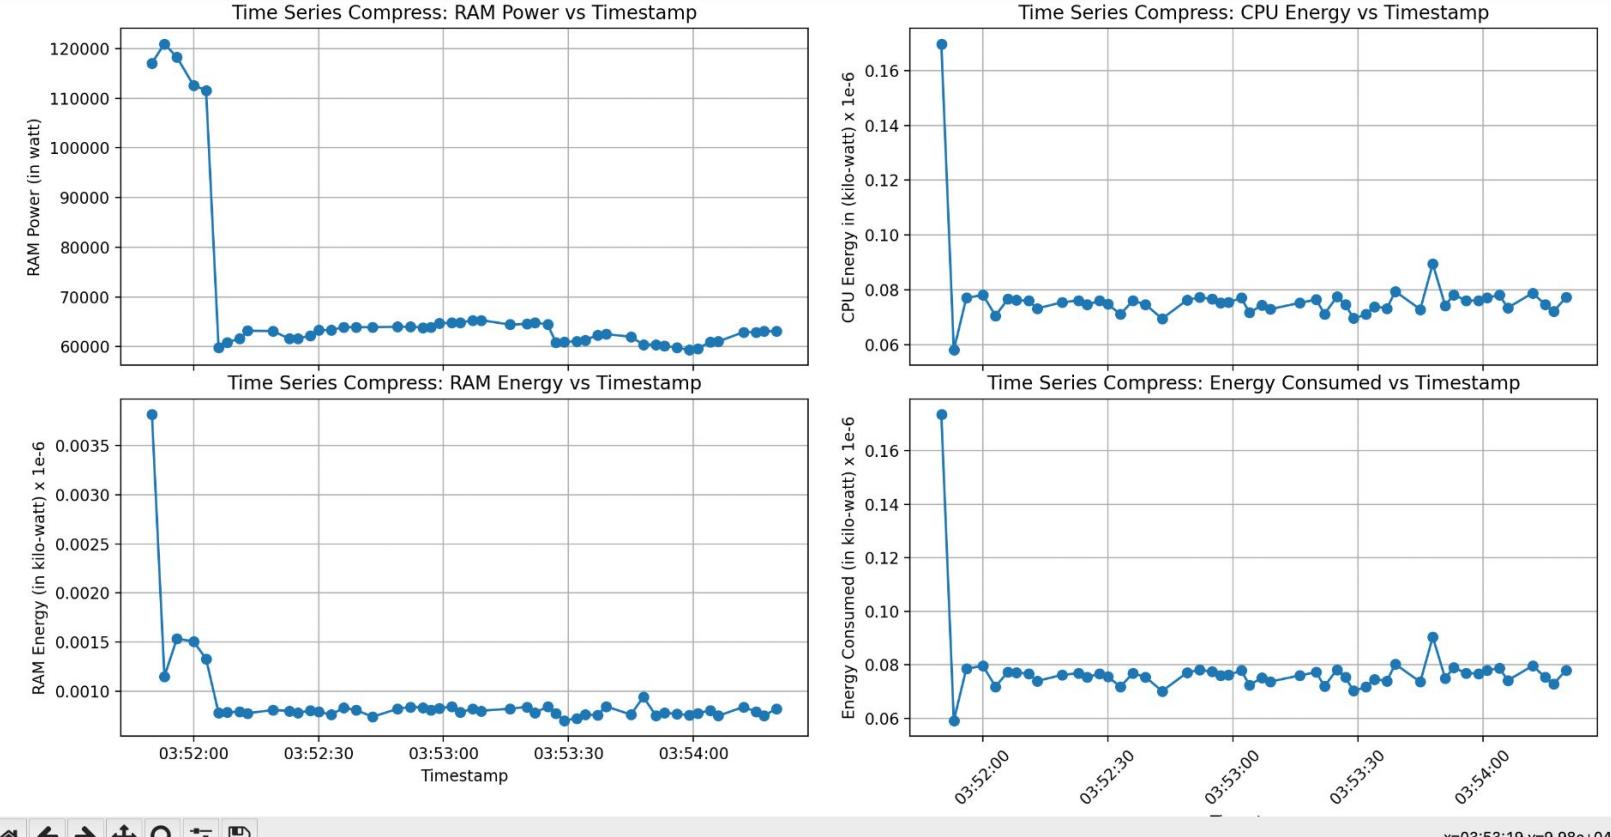


Figure 1

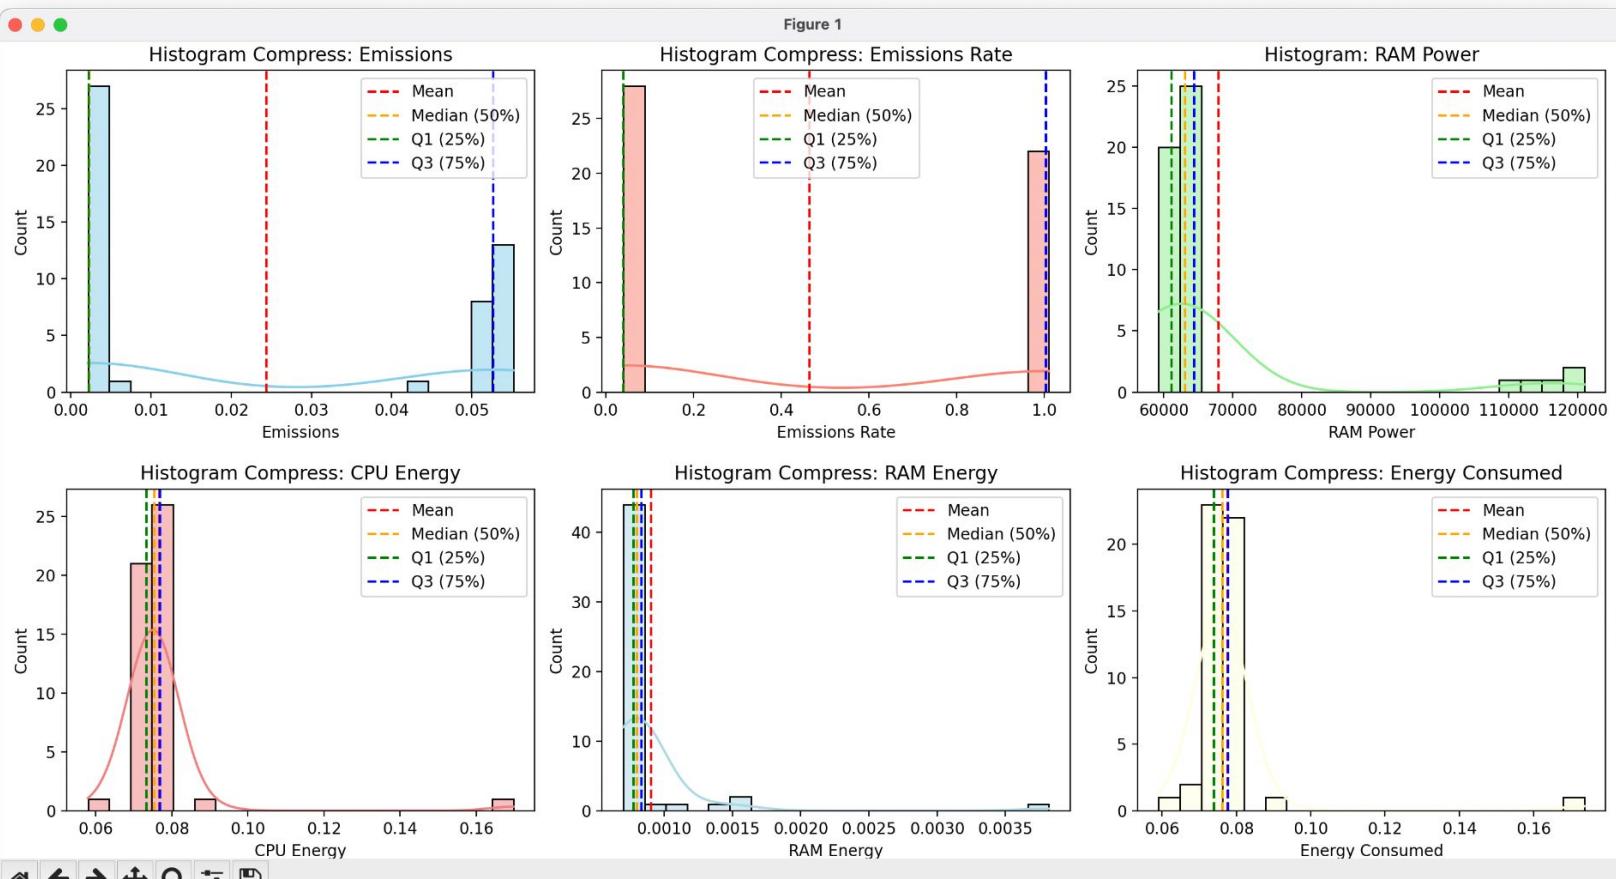


Figure 1

Time Series Decompress: Duration vs Timestamp

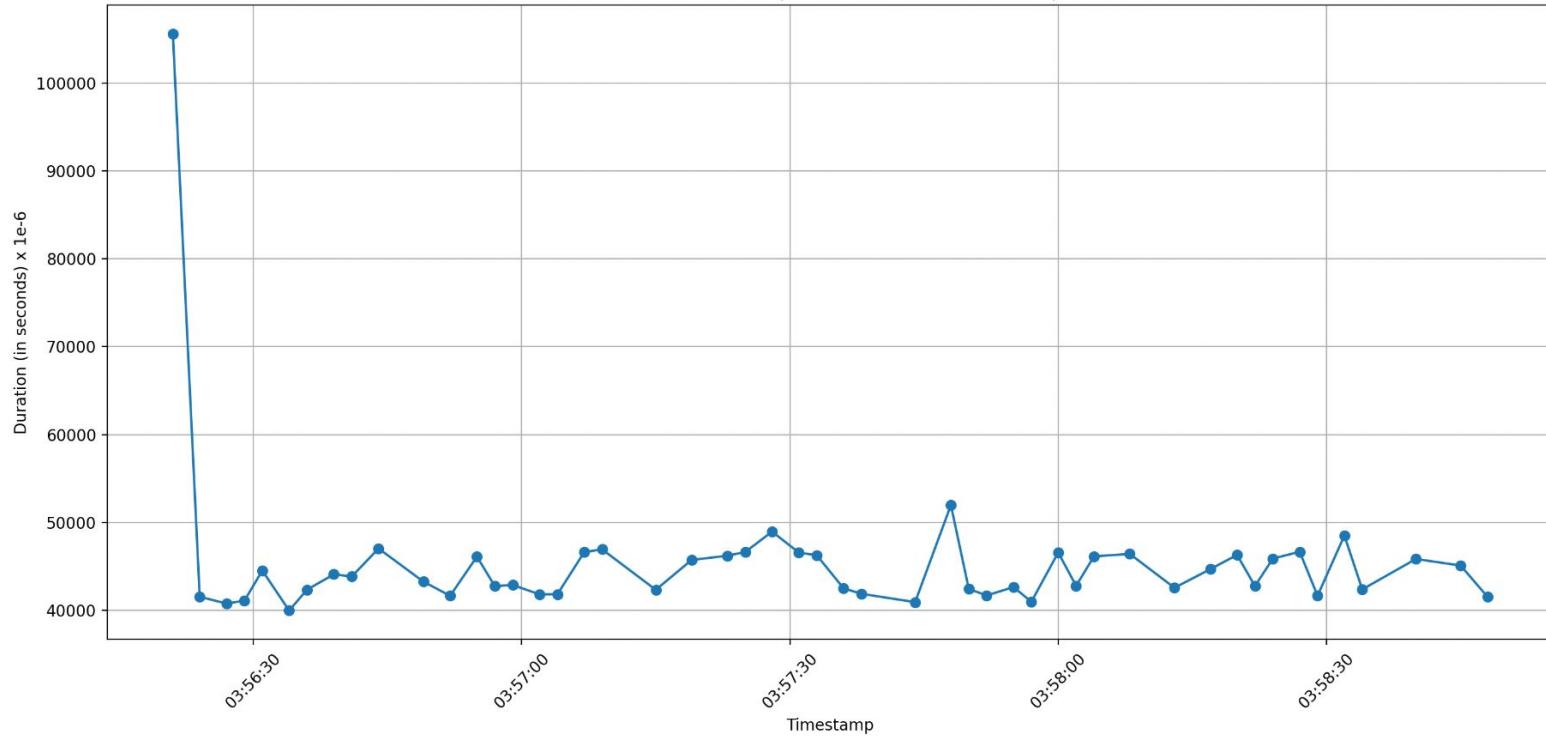
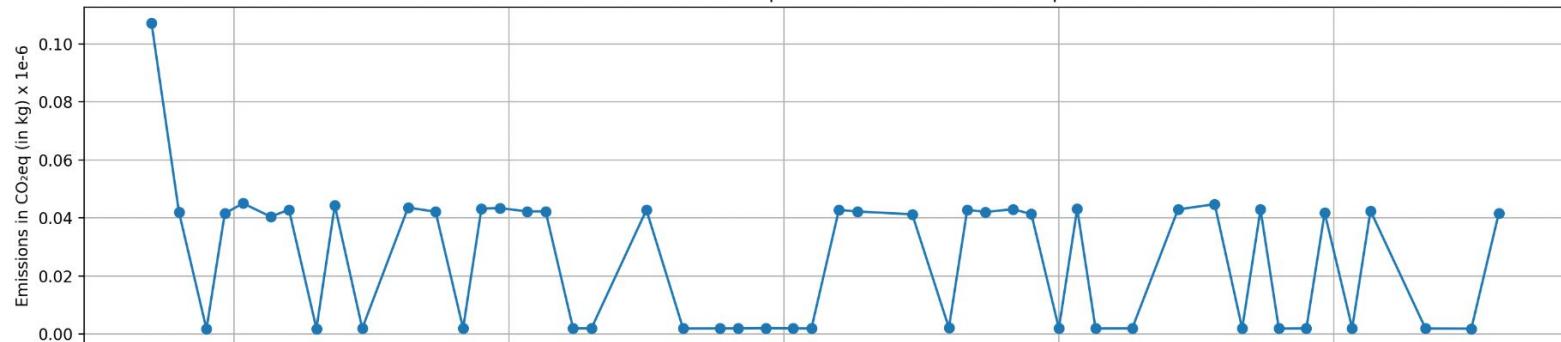
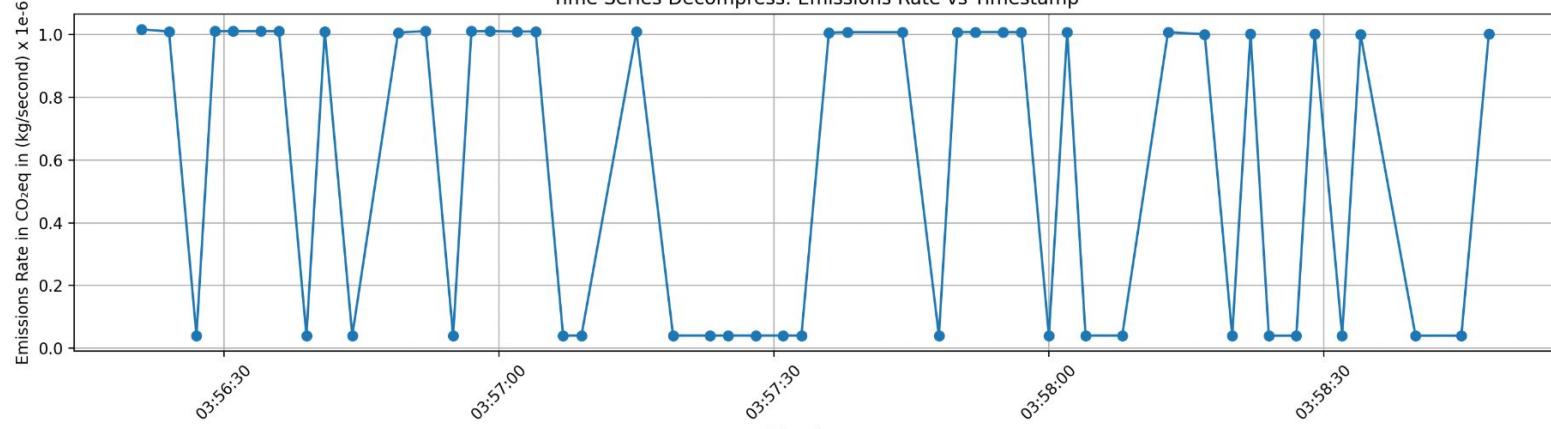


Figure 1

Time Series Decompress: Emissions vs Timestamp



Time Series Decompress: Emissions Rate vs Timestamp



x=03:56:31 y=0.0457

Figure 1

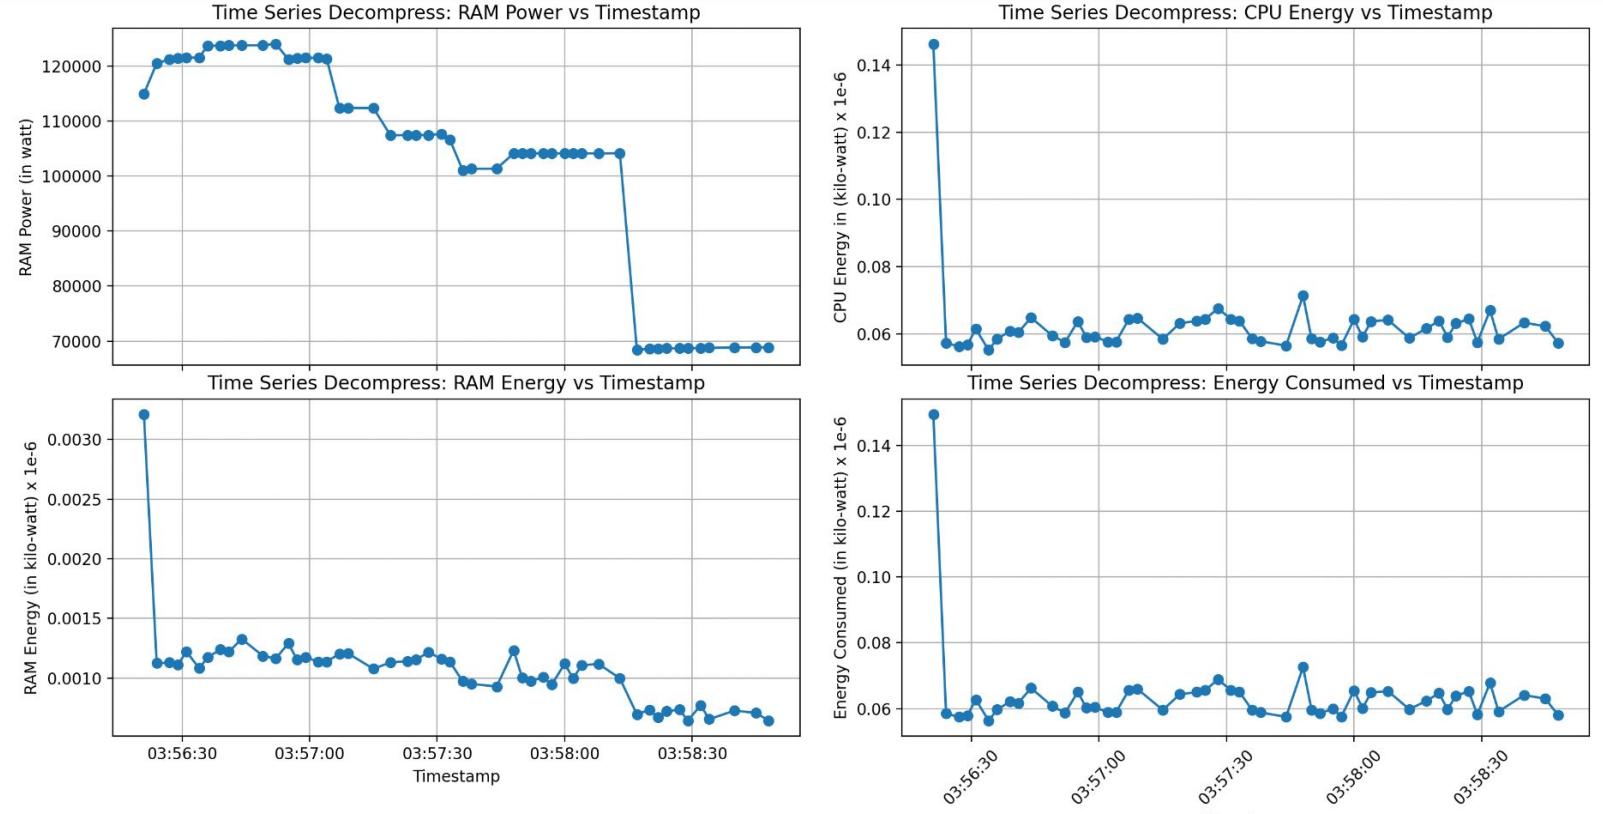
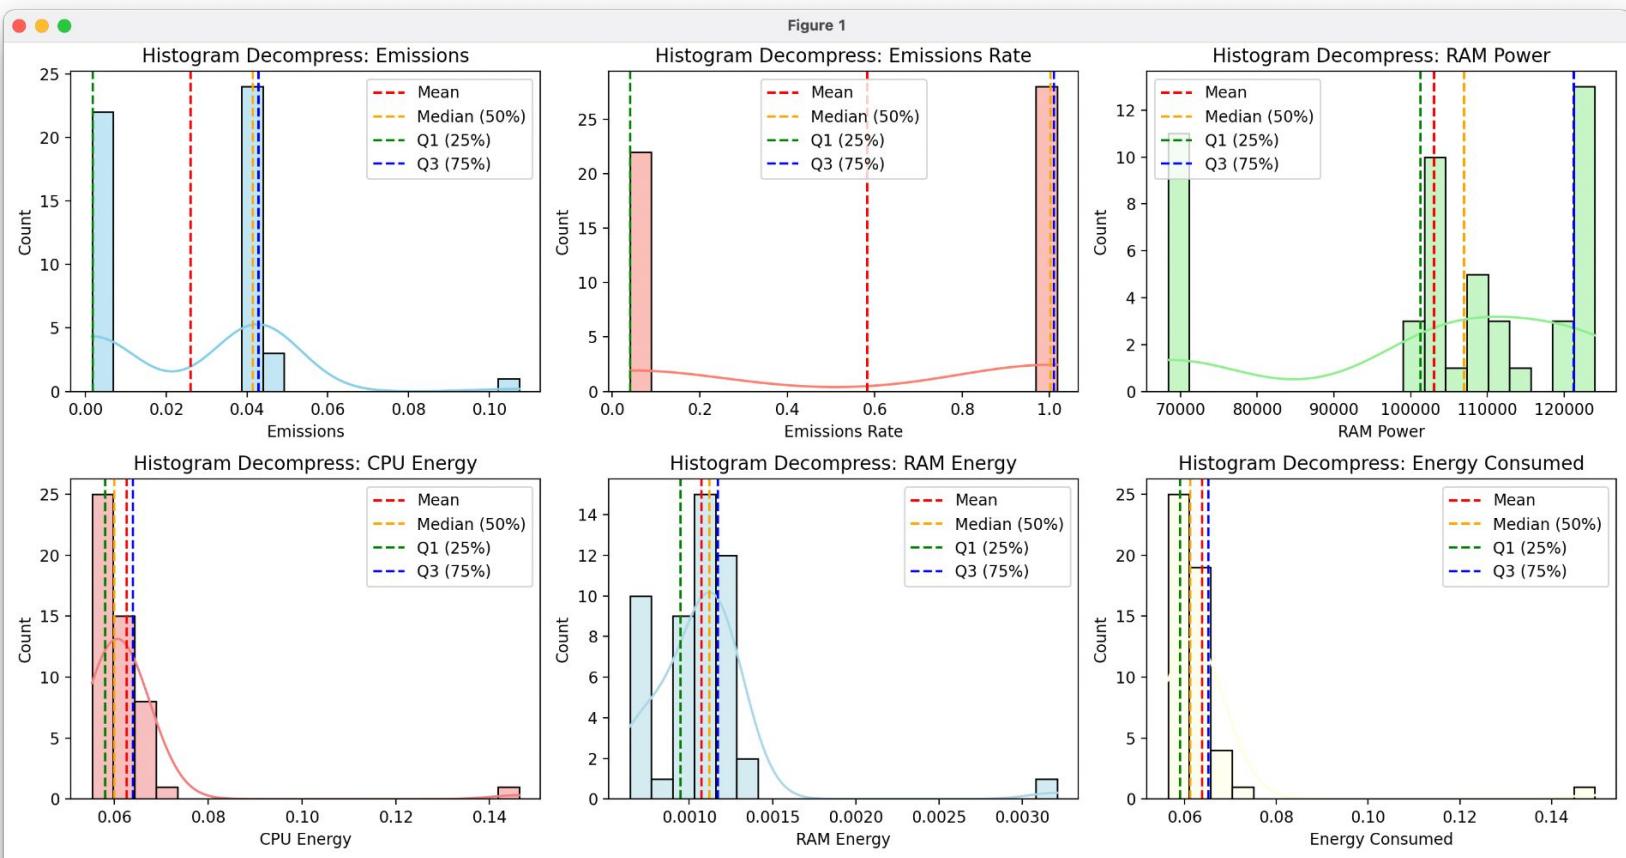


Figure 1



powermetrics and influxdb

Powermetrics is a built-in tool for macOS which monitors CPU usage and determines how much CPU time and CPU power is being allocated to different quality of service classes (QoS). By assigning a QoS to work, you indicate its importance, and the system prioritizes it and schedules it accordingly

Link to the MAN page -

<https://www.unix.com/man-page/osx/1/powermetrics/>

powermetrics and influxdb

Live Demonstration!

Next Plans

- Test the results with another memory profiler memray
- Try out scaphandre
- Read more about refactoring python code and try implementing some of them
- https://aman.ai/tnt/docs/Chapter5/speed_up_code.html
- <https://pythonspeed.com/pygotham19>
- Write blogs for Mid-Term Evaluation

Link to the GSoC Blogs – <https://manaspratimbiswas.com/gsoc>

Blockers

- Some dependencies and tools are not running on my current Ubuntu 22.04.2 LTS
- Need to downgrade and install Ubuntu 18.04.6
- Need intel machines with Nvidia to test out some tools